

CENTRAL TEST

March 17, 2024

```
[21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import style
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression
```

```
[22]: # Load the dataset
data = pd.read_csv("C:\\Users\\HP\\Desktop\\kiva_mpi_region_locations.csv")
data
```

```
[22]:
```

| | LocationName | ISO | country | region | world_region | \ |
|------|-------------------------|-----|-------------|------------|--------------|-----|
| 0 | Badakhshan, Afghanistan | AFG | Afghanistan | Badakhshan | South Asia | |
| 1 | Badghis, Afghanistan | AFG | Afghanistan | Badghis | South Asia | |
| 2 | Baghlan, Afghanistan | AFG | Afghanistan | Baghlan | South Asia | |
| 3 | Balkh, Afghanistan | AFG | Afghanistan | Balkh | South Asia | |
| 4 | Bamyan, Afghanistan | AFG | Afghanistan | Bamyan | South Asia | |
| ... | ... | ... | ... | ... | ... | |
| 2767 | | NaN | NaN | NaN | NaN | NaN |
| 2768 | | NaN | NaN | NaN | NaN | NaN |
| 2769 | | NaN | NaN | NaN | NaN | NaN |
| 2770 | | NaN | NaN | NaN | NaN | NaN |
| 2771 | | NaN | NaN | NaN | NaN | NaN |

| | MPI | geo | lat | lon |
|------|-------|---------------------------------|-----------|-----------|
| 0 | 0.387 | (36.7347725, 70.81199529999999) | 36.734772 | 70.811995 |
| 1 | 0.466 | (35.1671339, 63.7695384) | 35.167134 | 63.769538 |
| 2 | 0.300 | (35.8042947, 69.2877535) | 35.804295 | 69.287754 |
| 3 | 0.301 | (36.7550603, 66.8975372) | 36.755060 | 66.897537 |
| 4 | 0.325 | (34.8100067, 67.8212104) | 34.810007 | 67.821210 |
| ... | ... | ... | ... | ... |
| 2767 | NaN | (1000.0, 1000.0) | NaN | NaN |
| 2768 | NaN | (1000.0, 1000.0) | NaN | NaN |
| 2769 | NaN | (1000.0, 1000.0) | NaN | NaN |
| 2770 | NaN | (1000.0, 1000.0) | NaN | NaN |
| 2771 | NaN | (1000.0, 1000.0) | NaN | NaN |

[2772 rows x 9 columns]

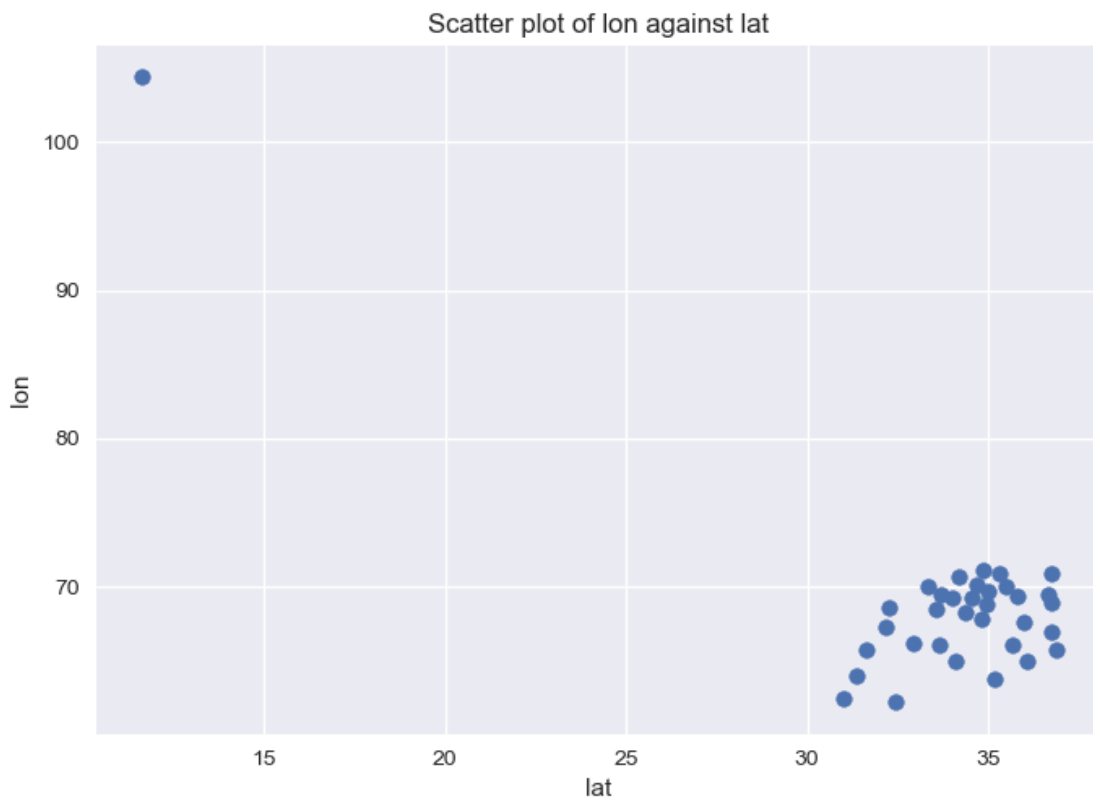
```
[23]: # Prepare the features and target variable
X = data[['lat']].iloc[:34] # Considering only the first 34 rows for simplicity
y = data['lon'].iloc[:34] # Assuming 'lon' as the target variable
```

```
[24]: # Plot a scatter graph
style.use('seaborn')
plt.xlabel("lat")
plt.ylabel("lon")
plt.scatter(X, y)
plt.title("Scatter plot of lon against lat")
plt.show()
```

C:\Users\HP\AppData\Local\Temp\ipykernel_9104\488471832.py:2:

MatplotlibDeprecationWarning: The seaborn styles shipped by Matplotlib are deprecated since 3.6, as they no longer correspond to the styles shipped by seaborn. However, they will remain available as 'seaborn-v0_8-*<style>*'. Alternatively, directly use the seaborn API instead.

```
style.use('seaborn')
```



```
[25]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳ random_state=42)
```

```
[26]: # Create a linear regression model
base_model = LinearRegression()
```

```
[27]: # Fit the base model
base_model.fit(X_train, y_train)
```

```
[27]: LinearRegression()
```

```
[28]: # Make predictions with the base model
base_pred = base_model.predict(X_test)
```

```
[29]: # Plot the regression line
plt.scatter(X_test, y_test, color='blue')
plt.plot(X_test, base_pred, color='red', linewidth=3)
plt.xlabel("lat")
plt.ylabel("lon")
plt.title("Linear Regression with base model")
plt.show()
```



```
[30]: # Calculate the coefficient and intercept of the base model
base_coef = base_model.coef_
```

```
base_intercept = base_model.intercept_  
print("Coefficient:", base_coef)  
print("Intercept:", base_intercept)
```

Coefficient: [-1.32226706]
Intercept: 113.00294245097427

```
[31]: # Calculate the R^2 score of the base model  
base_r2_score = base_model.score(X_test, y_test)  
print("R^2 Score of base model:", base_r2_score)
```

R² Score of base model: -1.8921327959431977

```
[32]: # Define parameter grid for grid search  
param_grid = {  
    'fit_intercept': [True, False],  
    'copy_X': [True, False],  
    'positive': [True, False]  
}
```

```
[33]: # Perform grid search  
grid_search = GridSearchCV(base_model, param_grid, cv=5)  
grid_search.fit(X_train, y_train)
```

```
[33]: GridSearchCV(cv=5, estimator=LinearRegression(),  
                  param_grid={'copy_X': [True, False],  
                              'fit_intercept': [True, False],  
                              'positive': [True, False]})
```

```
[34]: # Get the best parameters found by grid search  
best_params = grid_search.best_params_  
print("Best Parameters:", best_params)
```

Best Parameters: {'copy_X': True, 'fit_intercept': True, 'positive': True}

```
[35]: # Create a new model with the best parameters  
optimized_model = LinearRegression(**best_params)
```

```
[36]: # Fit the optimized model  
optimized_model.fit(X_train, y_train)
```

```
[36]: LinearRegression(positive=True)
```

```
[37]: # Make predictions with the optimized model  
optimized_pred = optimized_model.predict(X_test)
```

```
[38]: # Plot the regression line for optimized model  
plt.scatter(X_test, y_test, color='blue')  
plt.plot(X_test, optimized_pred, color='green', linewidth=3)
```

```
plt.xlabel("lat")
plt.ylabel("lon")
plt.title("Linear Regression with optimized model")
plt.show()
```



```
[39]: # Calculate the coefficient and intercept of the optimized model
optimized_coef = optimized_model.coef_
optimized_intercept = optimized_model.intercept_
print("Optimized Coefficient:", optimized_coef)
print("Optimized Intercept:", optimized_intercept)
```

```
Optimized Coefficient: [0.]
Optimized Intercept: 68.725456000000001
```

```
[40]: # Calculate the R^2 score of the optimized model
optimized_r2_score = optimized_model.score(X_test, y_test)
print("R^2 Score of optimized model:", optimized_r2_score)
```

```
R^2 Score of optimized model: -0.013815565615129088
```

LOGISTIC REGRESSION

```
[36]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import KBinsDiscretizer
from sklearn.metrics import accuracy_score
```

```
[37]: # Load the dataset
data = pd.read_csv("C:\\Users\\HP\\Desktop\\activitycalorieintensitystepsJoined_
↳ Extract 1.csv")
data
```

```
[37]:
```

| | Id | ActivityDate | TotalSteps | TotalDistance | TrackerDistance | \ |
|-----|------------|--------------|------------|---------------|-----------------|---|
| 0 | 1624580081 | 05-01-16 | 36019 | 28.030001 | 28.030001 | |
| 1 | 1644430081 | 4/14/2016 | 11037 | 8.020000 | 8.020000 | |
| 2 | 1644430081 | 4/19/2016 | 11256 | 8.180000 | 8.180000 | |
| 3 | 1644430081 | 4/28/2016 | 9405 | 6.840000 | 6.840000 | |
| 4 | 1644430081 | 4/30/2016 | 18213 | 13.240000 | 13.240000 | |
| .. | ... | ... | ... | ... | ... | |
| 935 | 1844505072 | 4/20/2016 | 8 | 0.010000 | 0.010000 | |
| 936 | 4020332650 | 4/17/2016 | 16 | 0.010000 | 0.010000 | |
| 937 | 4319703577 | 05-12-16 | 17 | 0.010000 | 0.010000 | |
| 938 | 6775888955 | 05-03-16 | 9 | 0.010000 | 0.010000 | |
| 939 | 7086361926 | 4/16/2016 | 31 | 0.010000 | 0.010000 | |

| | LoggedActivitiesDistance | VeryActiveDistance | ModeratelyActiveDistance | \ |
|-----|--------------------------|--------------------|--------------------------|---|
| 0 | 0.0 | 21.92 | 4.19 | |
| 1 | 0.0 | 0.36 | 2.56 | |
| 2 | 0.0 | 0.36 | 2.53 | |
| 3 | 0.0 | 0.20 | 2.32 | |
| 4 | 0.0 | 0.63 | 3.14 | |
| .. | ... | ... | ... | |
| 935 | 0.0 | 0.00 | 0.00 | |
| 936 | 0.0 | 0.00 | 0.00 | |
| 937 | 0.0 | 0.00 | 0.00 | |
| 938 | 0.0 | 0.00 | 0.00 | |
| 939 | 0.0 | 0.00 | 0.00 | |

| | LightActiveDistance | SedentaryActiveDistance | ... | \ |
|-----|---------------------|-------------------------|-----|---|
| 0 | 1.91 | 0.02 | ... | |
| 1 | 5.10 | 0.00 | ... | |
| 2 | 5.30 | 0.00 | ... | |
| 3 | 4.31 | 0.00 | ... | |
| 4 | 9.46 | 0.00 | ... | |
| .. | ... | ... | ... | |
| 935 | 0.01 | 0.00 | ... | |
| 936 | 0.01 | 0.00 | ... | |
| 937 | 0.01 | 0.00 | ... | |

| | | | |
|-----|------|------|-----|
| 938 | 0.01 | 0.00 | ... |
| 939 | 0.01 | 0.00 | ... |

| | LightlyActiveMinutes_1 | FairlyActiveMinutes_1 | VeryActiveMinutes_1 | \ |
|-----|------------------------|-----------------------|---------------------|---|
| 0 | 171 | 63 | 186 | |
| 1 | 252 | 58 | 5 | |
| 2 | 278 | 58 | 5 | |
| 3 | 227 | 53 | 3 | |
| 4 | 402 | 71 | 9 | |
| .. | ... | ... | ... | |
| 935 | 1 | 0 | 0 | |
| 936 | 2 | 0 | 0 | |
| 937 | 2 | 0 | 0 | |
| 938 | 1 | 0 | 0 | |
| 939 | 3 | 0 | 0 | |

| | SedentaryActiveDistance_1 | LightActiveDistance_1 | \ |
|-----|---------------------------|-----------------------|---|
| 0 | 0.02 | 1.91 | |
| 1 | 0.00 | 5.10 | |
| 2 | 0.00 | 5.30 | |
| 3 | 0.00 | 4.31 | |
| 4 | 0.00 | 9.46 | |
| .. | ... | ... | |
| 935 | 0.00 | 0.01 | |
| 936 | 0.00 | 0.01 | |
| 937 | 0.00 | 0.01 | |
| 938 | 0.00 | 0.01 | |
| 939 | 0.00 | 0.01 | |

| | ModeratelyActiveDistance_1 | VeryActiveDistance_1 | Id_1_1 | \ |
|-----|----------------------------|----------------------|------------|---|
| 0 | 4.19 | 21.92 | 1624580081 | |
| 1 | 2.56 | 0.36 | 1644430081 | |
| 2 | 2.53 | 0.36 | 1644430081 | |
| 3 | 2.32 | 0.20 | 1644430081 | |
| 4 | 3.14 | 0.63 | 1644430081 | |
| .. | ... | ... | ... | |
| 935 | 0.00 | 0.00 | 1844505072 | |
| 936 | 0.00 | 0.00 | 4020332650 | |
| 937 | 0.00 | 0.00 | 4319703577 | |
| 938 | 0.00 | 0.00 | 6775888955 | |
| 939 | 0.00 | 0.00 | 7086361926 | |

| | ActivityDay_1 | StepTotal |
|---|---------------|-----------|
| 0 | 05-01-16 | 36019 |
| 1 | 4/14/2016 | 11037 |
| 2 | 4/19/2016 | 11256 |
| 3 | 4/28/2016 | 9405 |

| | | |
|-----|-----------|-------|
| 4 | 4/30/2016 | 18213 |
| .. | ... | ... |
| 935 | 4/20/2016 | 8 |
| 936 | 4/17/2016 | 16 |
| 937 | 05-12-16 | 17 |
| 938 | 05-03-16 | 9 |
| 939 | 4/16/2016 | 31 |

[940 rows x 31 columns]

```
[38]: # Drop any non-numeric columns or columns causing errors (like date columns)
# For simplicity, we'll drop all non-numeric columns for now
numeric_data = data.select_dtypes(include='number')
```

```
[39]: # Split the data into features (X) and target variable (y)
X = numeric_data.drop(columns=['TotalDistance']) # Features
y = numeric_data['TotalDistance'] # Target variable
```

```
[40]: # Convert the target variable into categories or bins
# Here, we'll use KBinsDiscretizer to convert it into 5 bins
est = KBinsDiscretizer(n_bins=5, encode='ordinal', strategy='uniform')
y_bins = est.fit_transform(y.values.reshape(-1, 1)).astype(int).flatten()
```

C:\Users\HP\anaconda3\Lib\site-packages\sklearn\preprocessing_discretization.py:239: FutureWarning: In version 1.5 onwards, subsample=200_000 will be used by default. Set subsample explicitly to silence this warning in the mean time. Set subsample=None to disable subsampling explicitly.

warnings.warn(

```
[41]: # Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y_bins, test_size=0.2,
↳ random_state=42)
```

```
[42]: # Initialize the logistic regression model
model = LogisticRegression(max_iter=1000)
```

```
[44]: # Train the model on the training data
model.fit(X_train, y_train)
```

```
[44]: LogisticRegression(max_iter=1000)
```

```
[51]: # Predict on the testing data
y_pred = model.predict(X_test)
```

```
[52]: # Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```


Accuracy: 0.5531914893617021

MODEL OPTIMIZATION

```
[53]: from sklearn.preprocessing import StandardScaler
      from sklearn.model_selection import GridSearchCV
```

```
[54]: # Feature Scaling
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

```
[55]: # Hyperparameter Tuning
      param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100]}
      grid_search = GridSearchCV(LogisticRegression(max_iter=1000), param_grid, cv=5)
      grid_search.fit(X_train_scaled, y_train)
```

C:\Users\HP\anaconda3\Lib\site-packages\sklearn\model_selection_split.py:725:
UserWarning: The least populated class in y has only 3 members, which is less
than n_splits=5.

warnings.warn(
C:\Users\HP\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:460:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
regression

```
    n_iter_i = _check_optimize_result(
```

```
[55]: GridSearchCV(cv=5, estimator=LogisticRegression(max_iter=1000),
              param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100]})
```

```
[56]: # Get the best model
      best_model = grid_search.best_estimator_
```

```
[57]: # Evaluate the best model
      accuracy = best_model.score(X_test_scaled, y_test)
      print("Accuracy:", accuracy)
      print("Best hyperparameters:", grid_search.best_params_)
```

Accuracy: 0.9627659574468085

Best hyperparameters: {'C': 100}